



Presto: SQL-on-Anything

DataWorks Summit, San Jose 2017

Martin Traverso, Facebook

Matt Fuller, Teradata



Let's Make Some Noise!

@DataWorksSummit

#DWS17

#prestodb

#facebook

#teradata

What is Presto?

- Open source distributed SQL query engine
 - Originally developed by Facebook
 - ANSI SQL compliant
 - Like Hive, it's not a database
- Key Differentiators
 - Performance & Scale
 - Cross platform query capability, not only SQL on Hadoop
 - Supports federated queries
 - Used in production at many well known web-scale companies
- Distributed under the Apache License, hosted on GitHub

zuora

YAHOO!
JAPAN

JD. 京东
.COM

FreeWheel

amazon

mercado
Libre.com

LinkedIn

jampp

Marin
SOFTWARE

looker

Comcast

WIX.com

TREASURE
DATA

airbnb



Walmart

Alibaba Group



UBER

slack

Bloomberg

GROUPON

GREE

COGO
labs

FINRA

twitter

trulia

shopify

Atlassian

AdRoll

SHAZAM

Pinterest

openspan

NASDAQ

Dropbox

CUEBIQ

NETFLIX

Presto in Action

facebook.

NETFLIX

UBER

twitter

GROUPON

airbnb

Dropbox

WIX

Atlassian

FINRA

Bloomberg

Linked

Multiple clusters (1000s of nodes total)
300PB in HDFS, MySQL, and Raptor
1000s users, 10-100s concurrent queries

Presto in Action

facebook

NETFLIX

UBER

twitter

GROUPON

250+ nodes on AWS
40+ PB stored in S3 (Parquet)
Over 650 users with 6K+ queries
daily

Bloomberg

airbnb

Linked

Dropbox

WIX

Atlassian

FINRA

Presto in Action

facebook

NETFLIX

UBER

twitter

GROUPON

airbnb

300+ nodes (2 dedicated clusters)
100K+ & 20K+ queries daily

Bloomberg

Linked

Dropbox

WIX

Atlassian

FINRA

Presto in Action

facebook

NETFLIX

UBER

twitter

GROUPON

airbnb

200+ nodes on-premises
Parquet nested data

Bloomberg

Linked

Dropbox

WIX

Atlassian

FINRA

Presto in Action

facebook

NETFLIX

UBER

twitter

GROUPON

airbnb

Dropbox

WIX

Atlassian

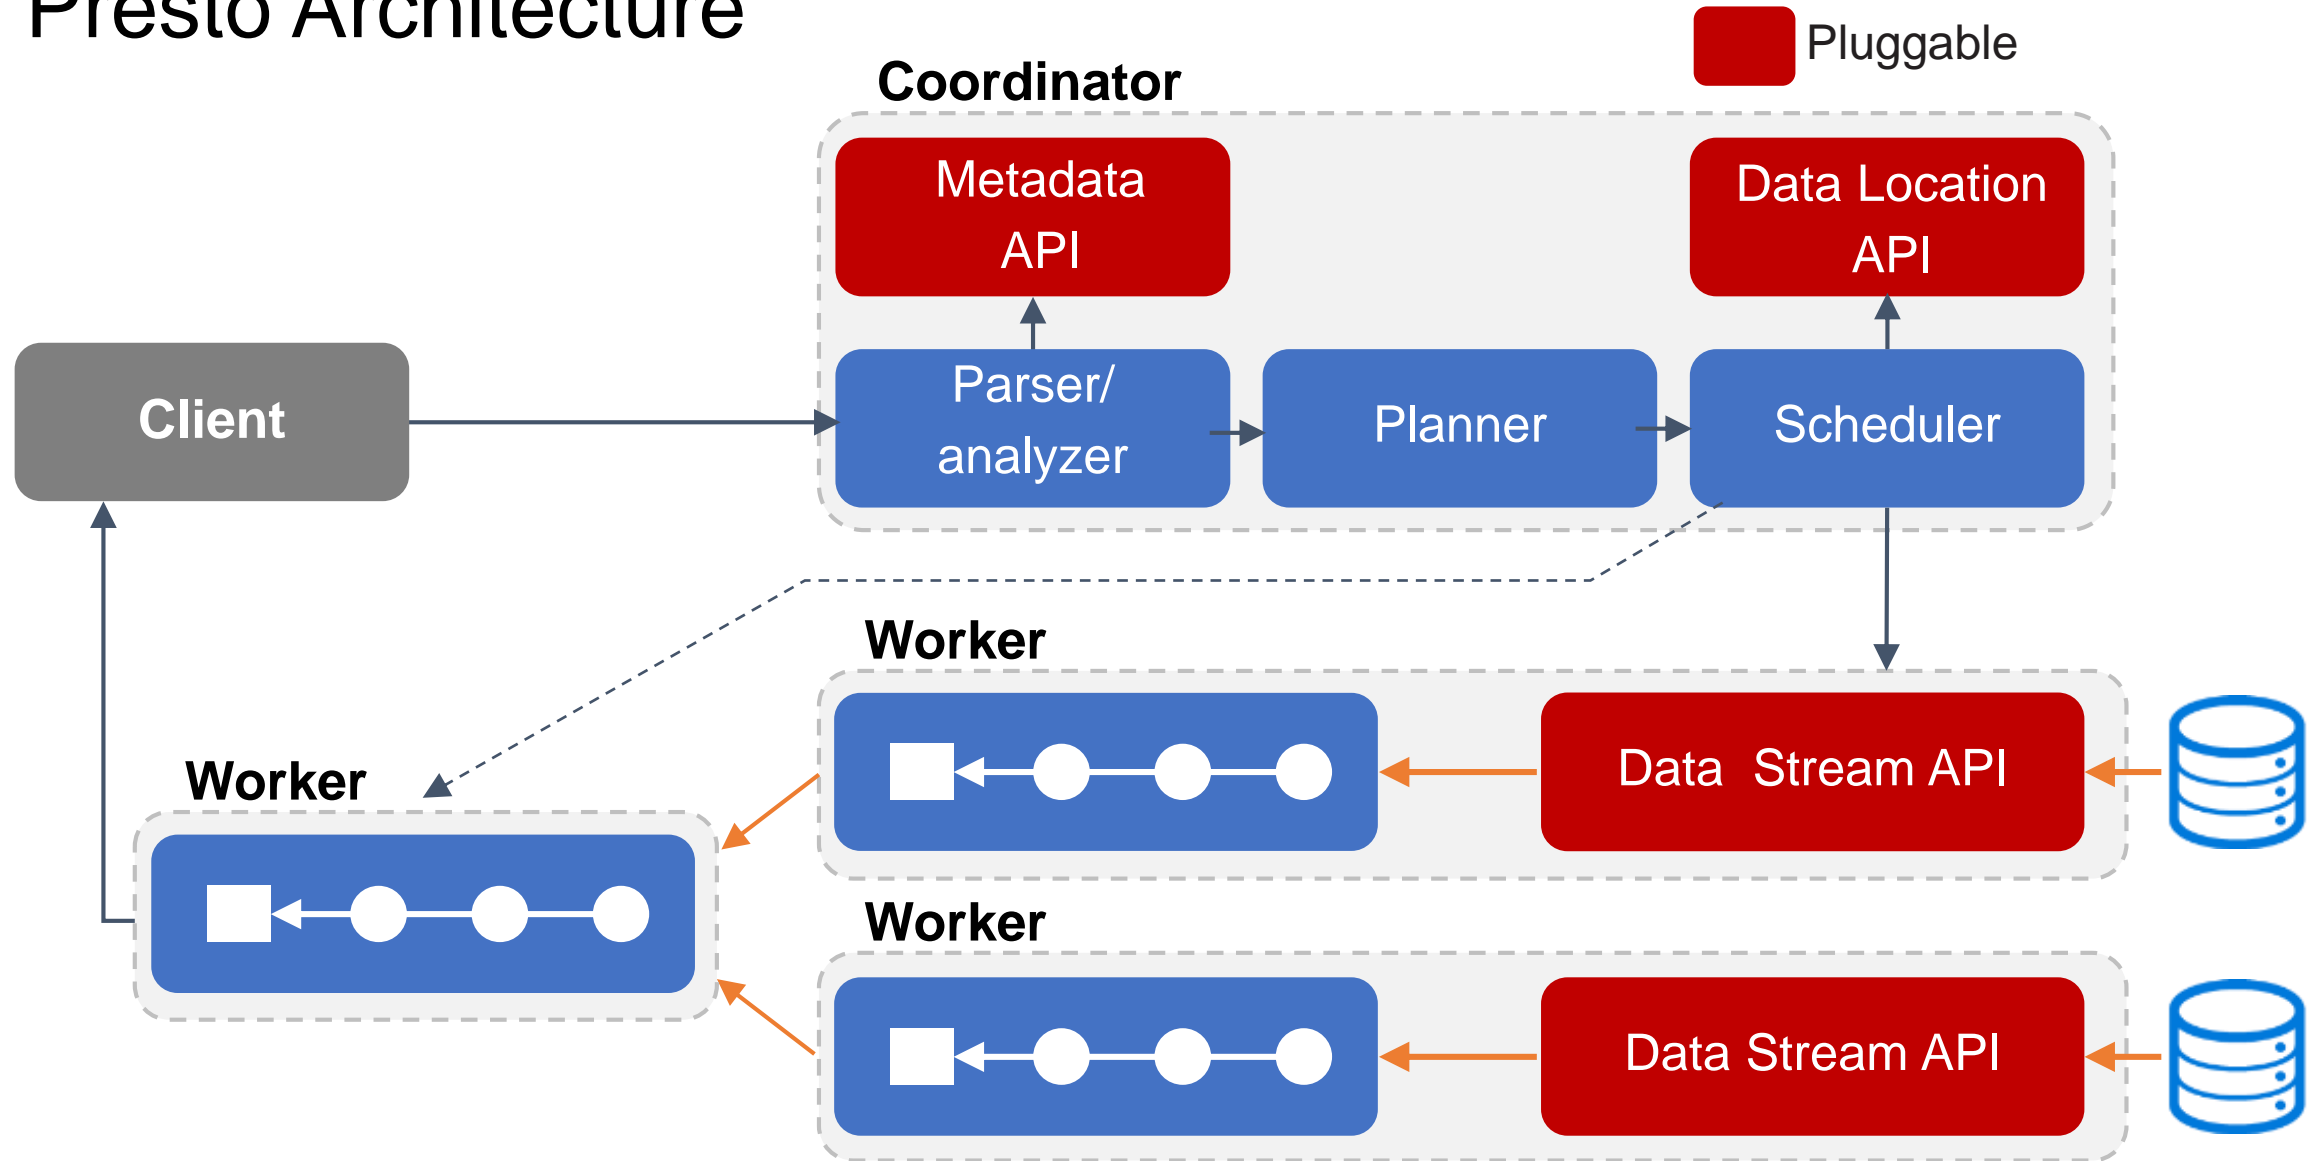
Bloomberg

LinkedIn

FINRA

120+ nodes in AWS
2PB is S3 and 200+ users
supported by Teradata

Presto Architecture

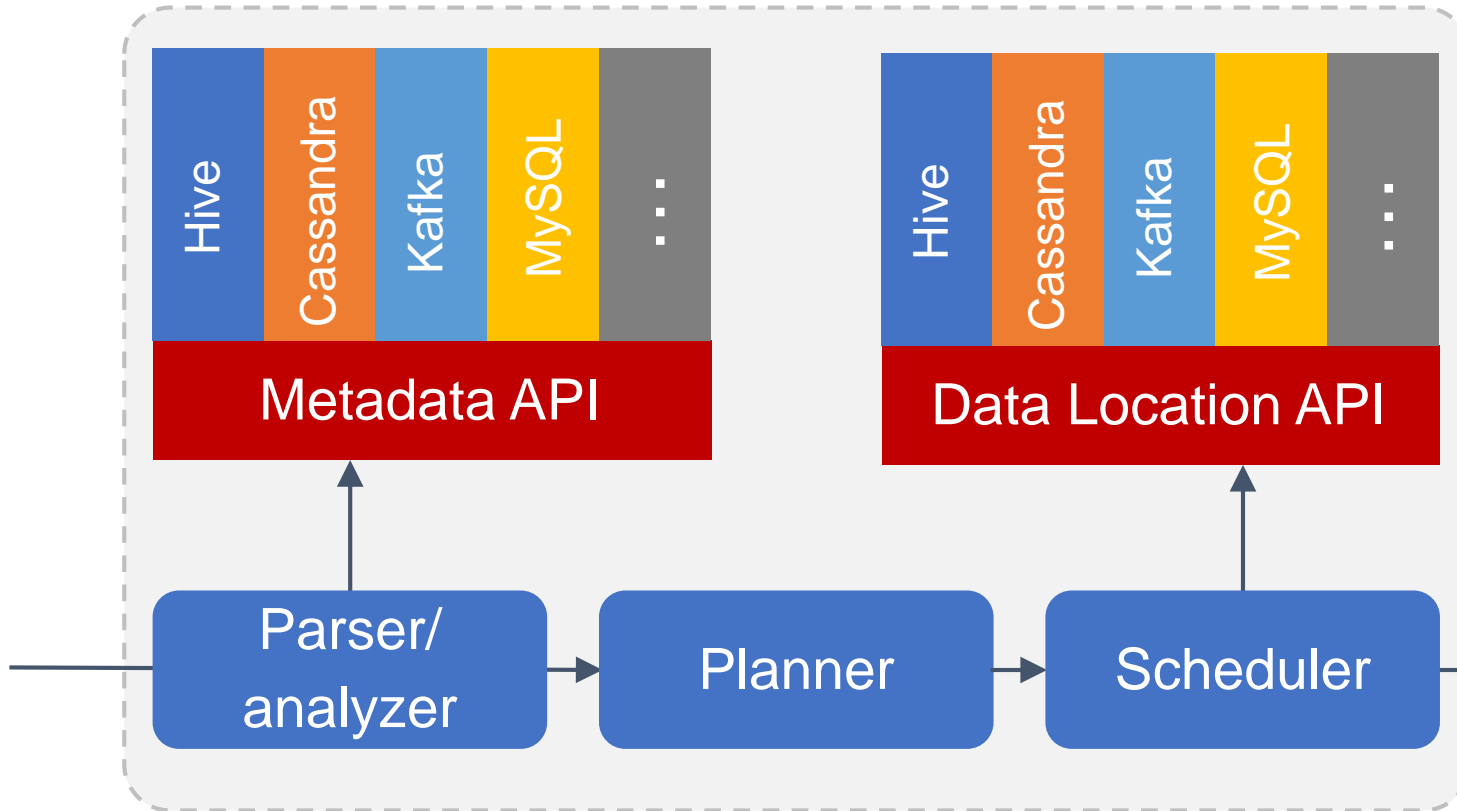


Presto is not a Database!

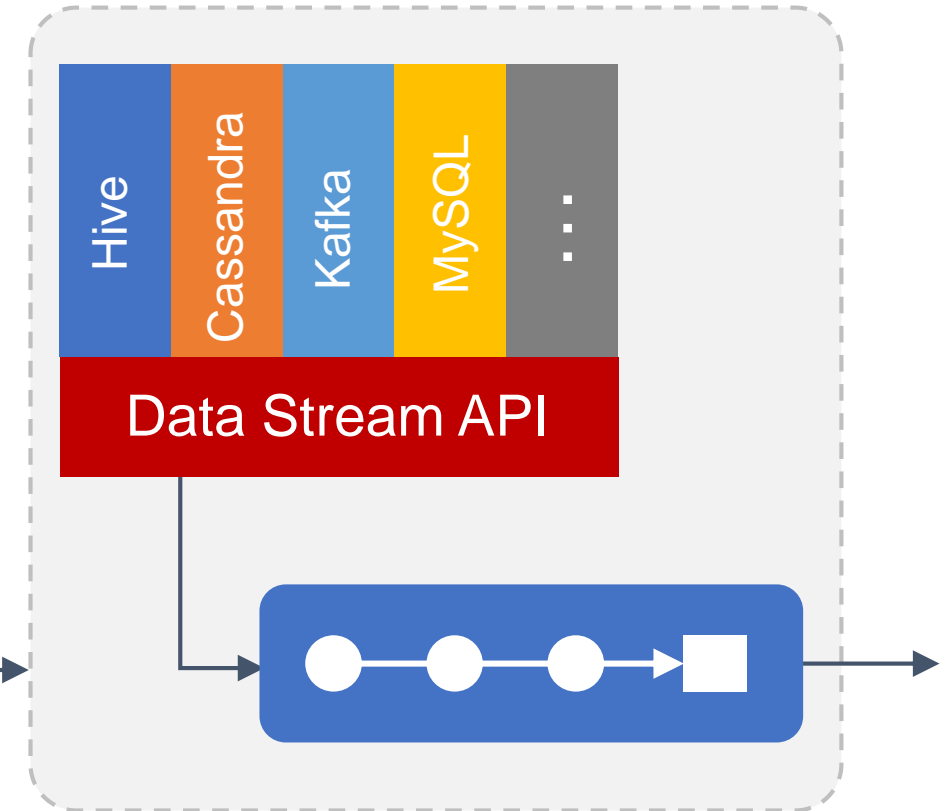
- Presto is a distributed query execution engine
 - Storage Independent
- Pluggable extensions
 - Connectors
 - Functions
 - Types
 - System access controllers
 - Resource group configuration managers
 - Event listeners
 - ...
- Built-in core functionalities
 - parser, execution, types, sql functions, monitoring

Presto Extensibility - Connector

Coordinator



Worker



Presto Connectors



PostgreSQL



presto



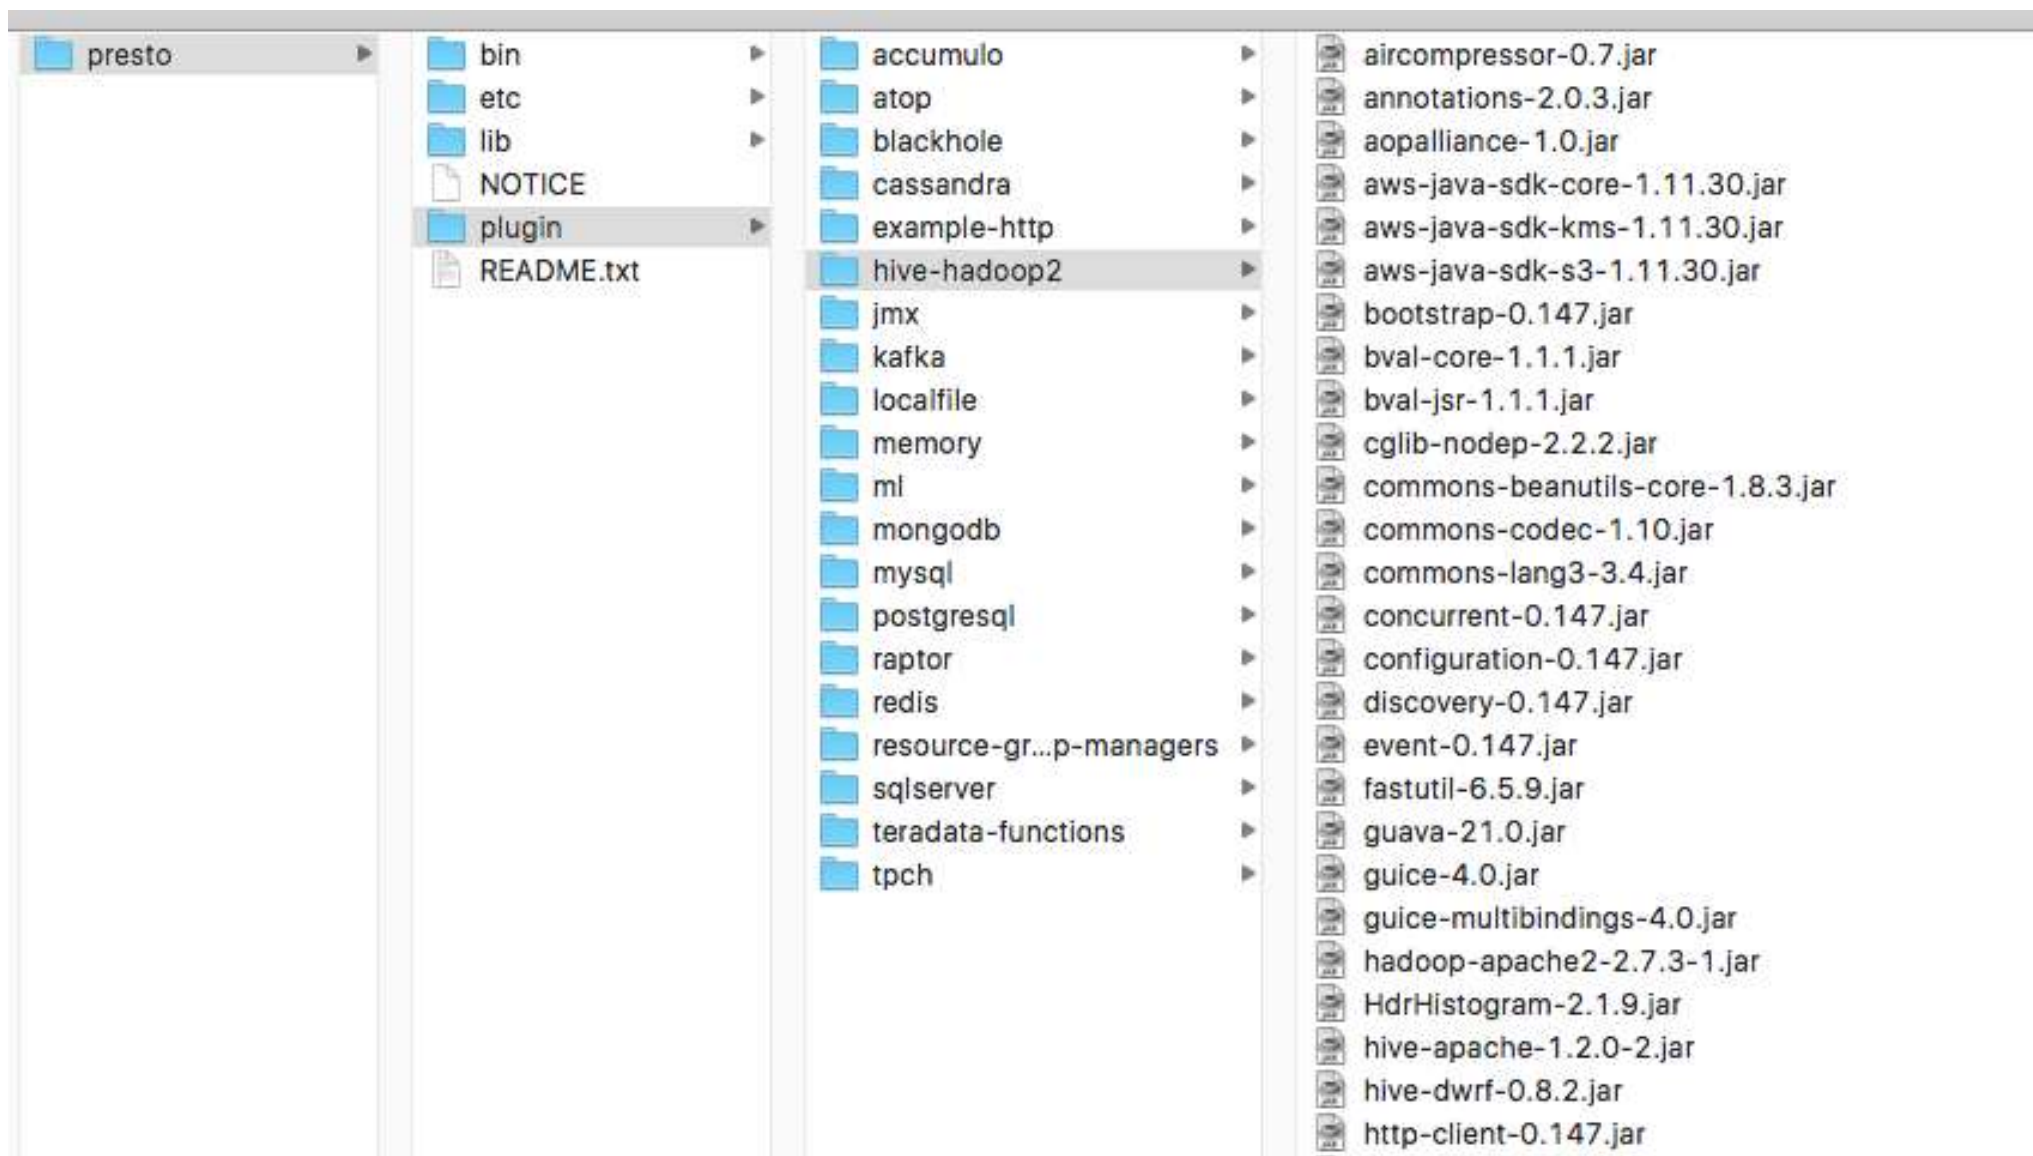
Amazon S3



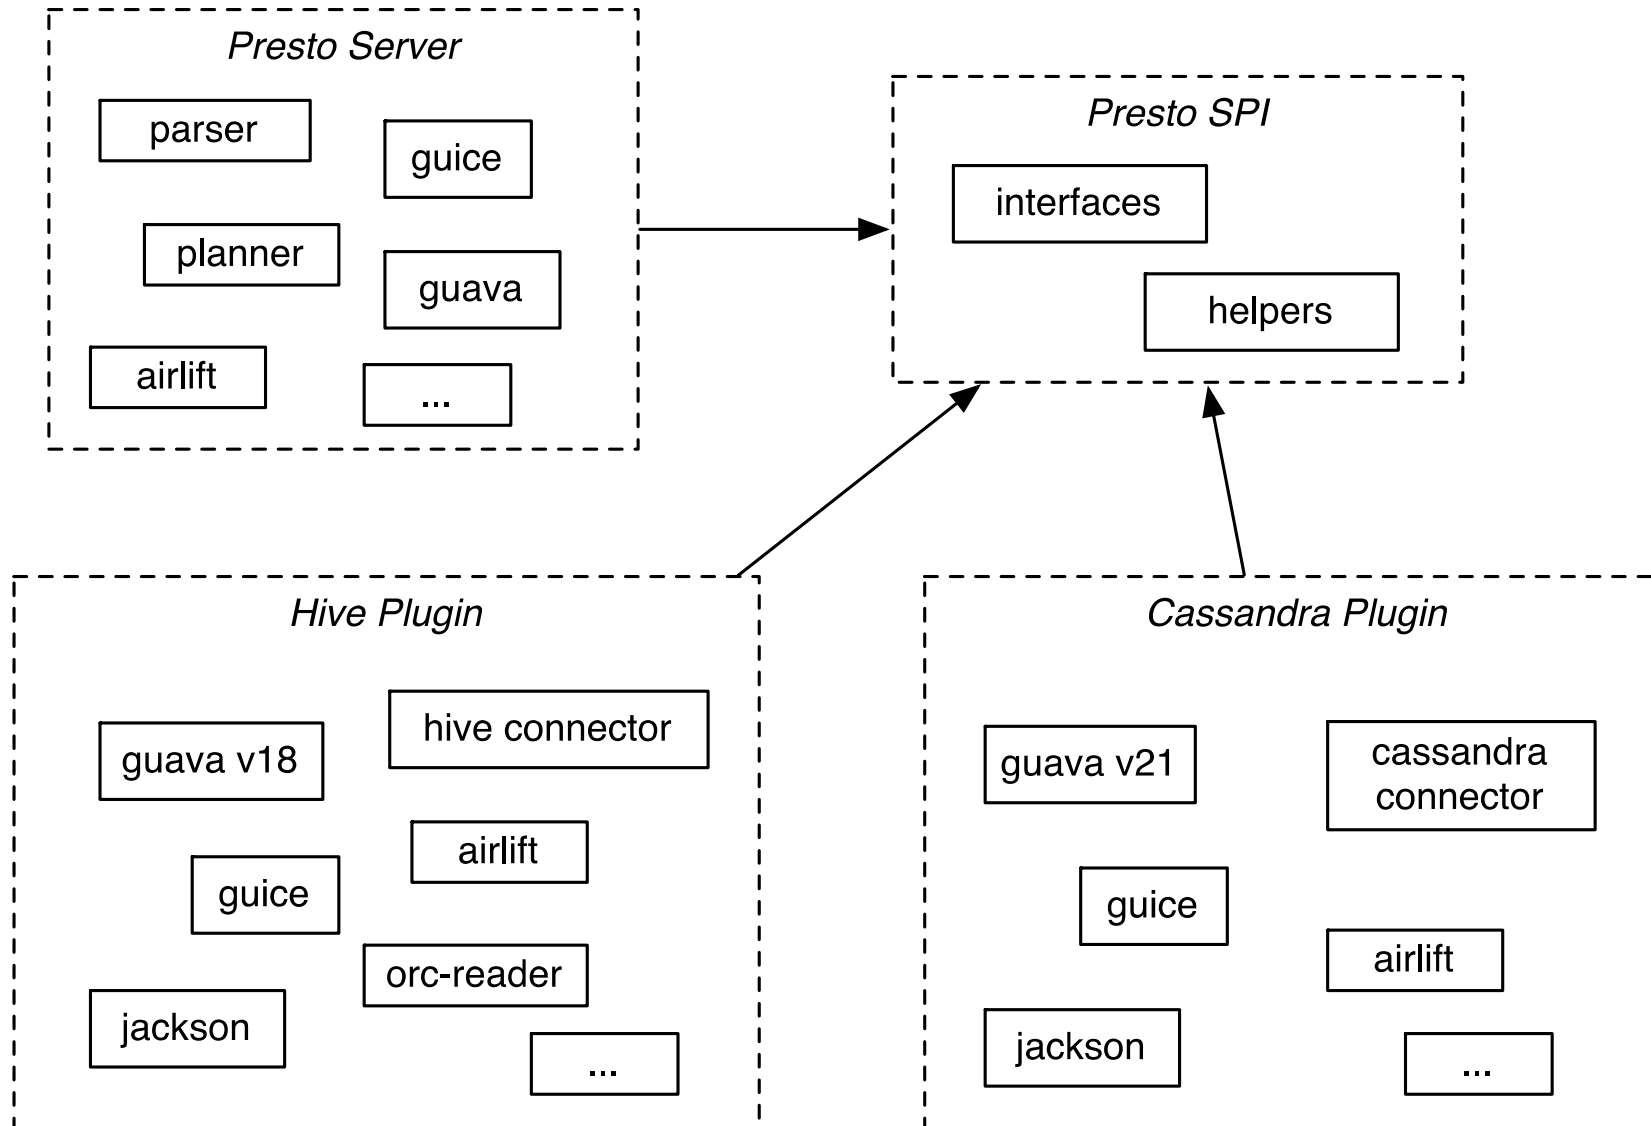
kafka



Plugins



Classloader Isolation



Plugin Interface

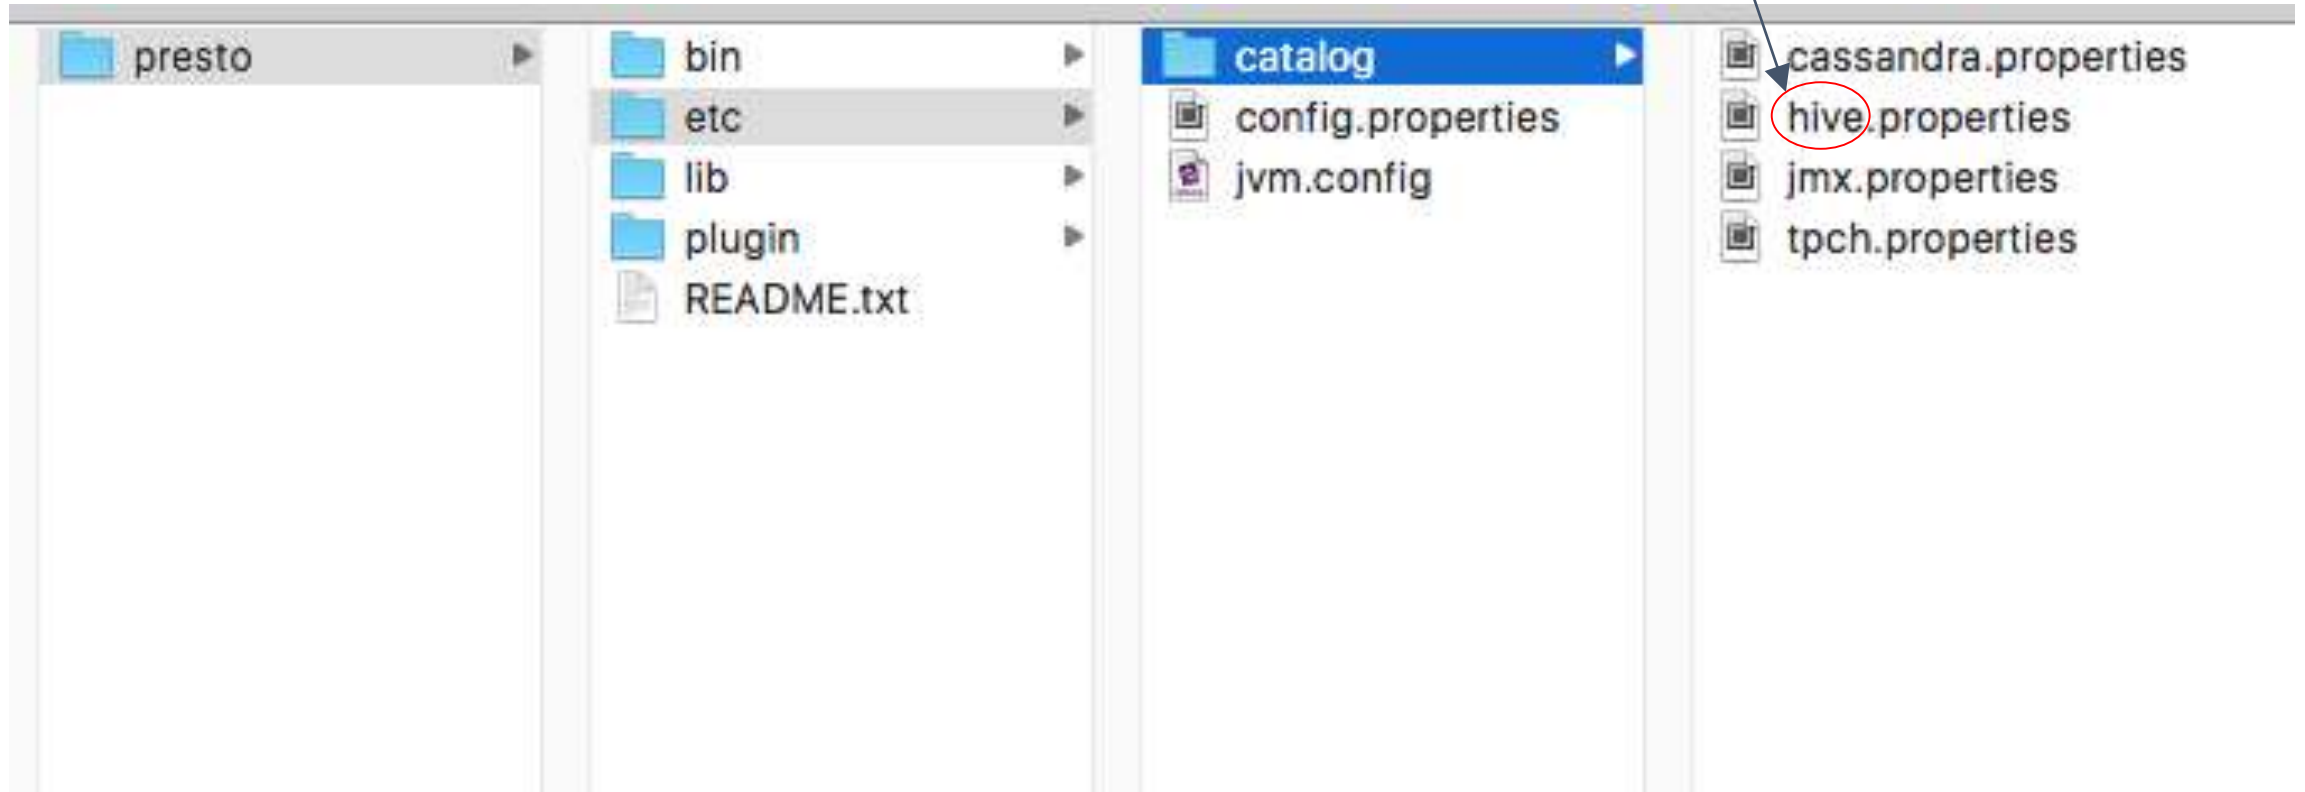
```
public interface Plugin
{
    default Iterable<ConnectorFactory> getConnectorFactories() { return emptyList(); }

    default Iterable<Type> getTypes() { return emptyList(); }
    default Set<Class<?>> getFunctions() { return emptySet(); }
    default Iterable<EventListenerFactory> getEventListenerFactories() { return emptyList(); }
    ...
}
```


Connector Configuration

- Catalog namespace owned by connector

Catalog Name



Connector Configuration

```
connector.name=hive-hadoop2
```

```
hive.metastore.uri=thrift://hadoop-master:9083
```

```
hive.metastore.authentication.type=KERBEROS
```

```
hive.metastore.service.principal=hive/hadoop-master@EXAMPLE.COM
```

```
hive.metastore.client.principal=hive/_HOST@EXAMPLE.COM
```

```
hive.metastore.client.keytab=/etc/presto/conf/hive-presto-master.keytab
```

```
hive.hdfs.authentication.type=KERBEROS
```

```
hive.hdfs.impersonation.enabled=true
```

```
hive.hdfs.presto.principal=presto-server/_HOST@EXAMPLE.COM
```

```
hive.hdfs.presto.keytab=/etc/presto/conf/presto-server.keytab
```

Query Analysis

```
SELECT custkey, count(*)  
FROM hive.tpch.orders  
WHERE orderstatus = '0'  
GROUP BY custkey
```

getTableHandle(...)

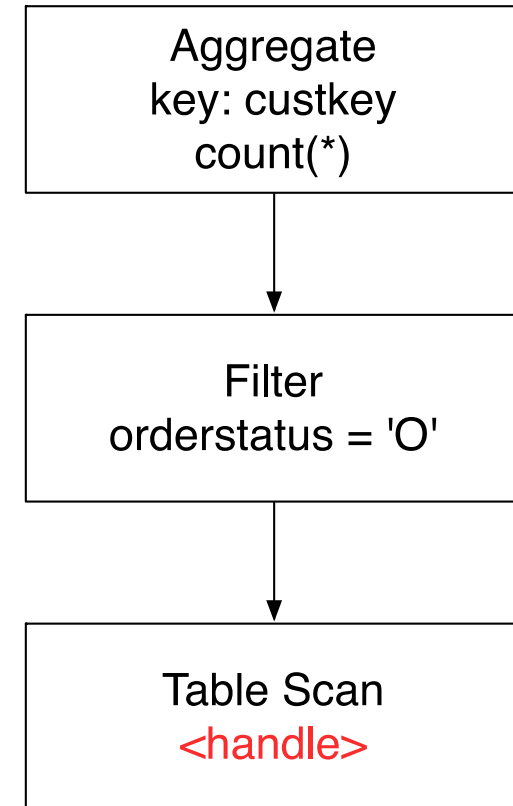
Table Handle

getTableMetadata(...)


```
orderkey BIGINT,  
custkey BIGINT,  
orderstatus VARCHAR(1),  
...
```

Query Planning

```
SELECT custkey, count(*)  
FROM hive.tpch.orders  
WHERE orderstatus = 'O'  
GROUP BY custkey
```



Query Execution




Split

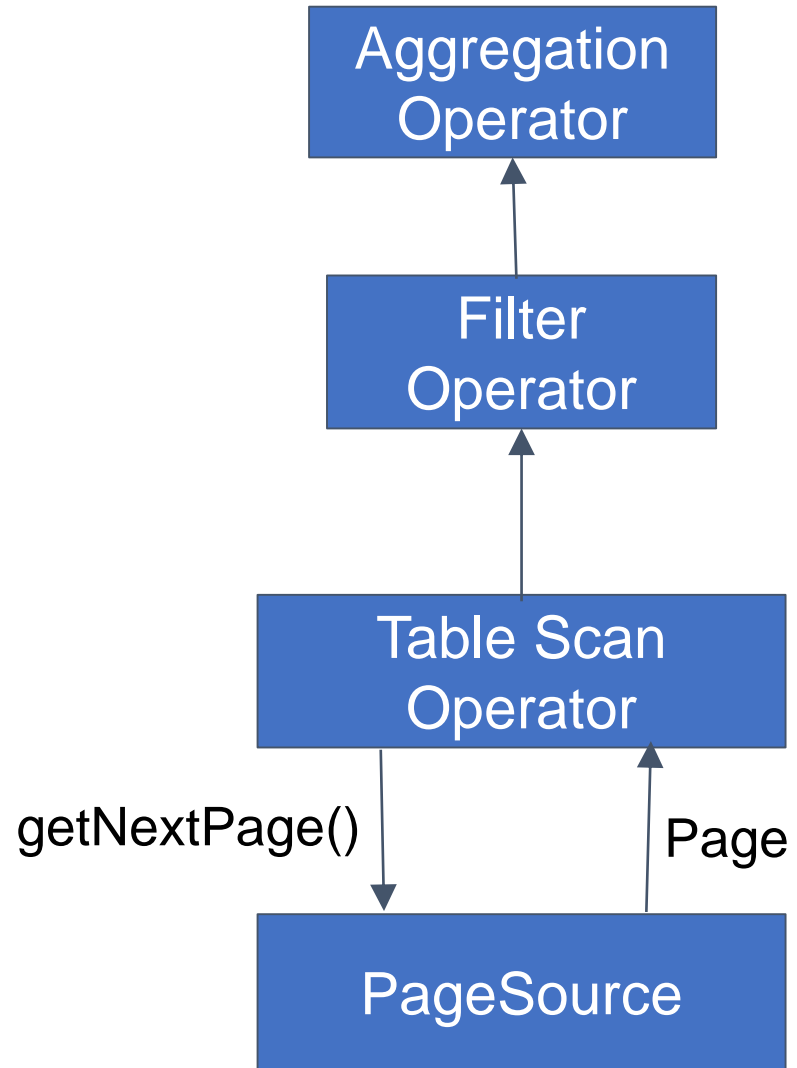
- Handle to logical chunk of a table
- Attributes
 - Remote access?
 - Location



Query Execution



Query Execution



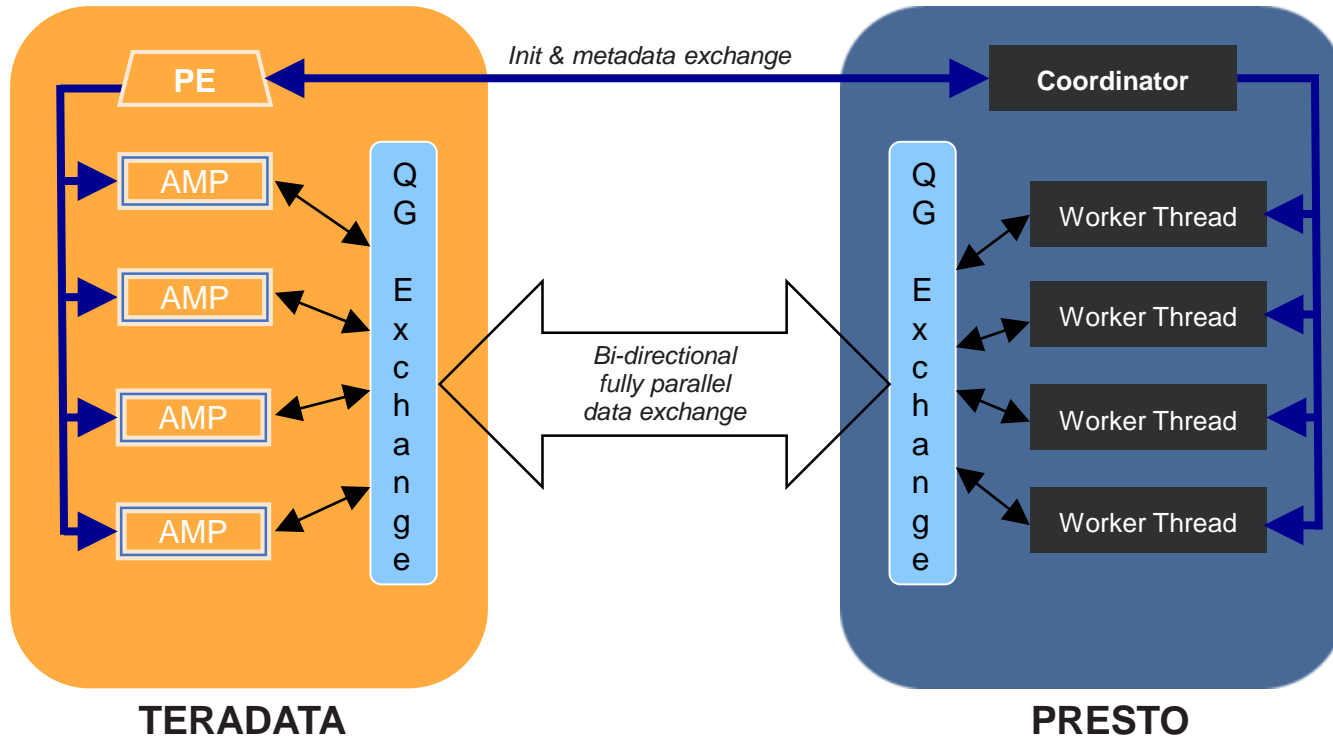
Presto Connectors @Facebook

- Hive connector
 - Warehouse (ad-hoc / batch)
- Raptor connector
 - Dashboards
 - Reporting backend for A/B testing framework
- Sharded MySQL connector
 - Reporting backend for user-facing products
- Other custom connectors for specialized data stores

Presto Connectors @Teradata Customers

- Teradata QueryGrid + Presto
 - Teradata, Hadoop, S3, Cassandra, RESTful
- Customer Use Cases
 - Recent sales data in **Teradata** needs to be joined with archived sales data that resided in **Hadoop**
 - Hadoop user using **Presto** needs to access pre-computed financial record in **Teradata**
 - Existing supplier data that is in Teradata is joined with archived product data that resides in **Amazon S3**

Teradata QueryGrid (powered by Presto)



- **Key features:**

- Low latency
- High performance
- Concurrency
- Pushdown
- Data conversion
- Compression
- Efficient CPU usage

Teradata QueryGrid SQL Examples

Teradata query joining data from Hadoop via Presto:

```
SELECT * FROM websales_current UNION ALL SELECT * FROM  
websales_archive @presto;
```

Presto query joining data in Teradata:

```
SELECT * FROM td.sales.websales current UNION ALL SELECT  
* FROM hive.sales.websales archive;
```

Conclusions

- Presto Connector API is expressive
- 3rd Party data source is 1st class citizen
- Single ANSI SQL to rule them all
 - Use BI tools on data which is not BI friendly
- Rapid data integration

Write your own connector!

- Issue SQL to GitHub!
 - <https://developer.github.com/v3/>
 - `SELECT count(*) FROM prestodb.presto.stargazers;`
- Connector Example
 - <https://github.com/prestodb/presto/tree/master/presto-example-http>
- Documentation
 - <https://prestodb.io/docs/current/develop.html>

Additional Resources

- Website
 - www.prestodb.io
- Presto Users Groups
 - www.groups.google.com/group/presto-users
- GitHub:
 - www.github.com/prestodb/presto
 - www.github.com/Teradata/presto (Teradata's development "fork")

